



# **AGENTIC AI:** FROM BUSINESS NEEDS TO WORKING AGENTS

**Professor Peyman Teymoori**



**BUSINESS  
NEEDS**

**WORKING AGENTS**



# Learning Objectives

1

## DEFINE

What makes software "agentic" and identify appropriate use cases

3

## TRANSLATE

A business problem into **measurable KPIs** and constraints

5

## CREATE

Specific, **testable success criteria** for agent behavior

2

## ARTICULATE

The **perceive → plan → act loop** and explain when it breaks down

4

## DESIGN

System instructions that encode **business logic and guardrails**

6

## IDENTIFY

Common risks (hallucination, cost, privacy) and **mitigation strategies**

# What Makes Software "Agentic"?



## Traditional Software

Input → Fixed Logic → Output

- Follows predetermined rules
- Executes exact commands
- No decision making

VS



## Agentic Software

Goal → Perceive → Plan → Act → Evaluate

- Pursues goals rather than commands
- Perceives its environment
- Makes decisions autonomously
- Uses tools to affect environment
- Learns from feedback within session



# The Core Agent Loop





# Components: Tools

## 💡 What are Tools?

Tools are **capabilities** the agent can invoke to interact with the world. They're the agent's hands and eyes.

### search\_database(query)

Returns records from database

### send\_email(to, subject, body)

Sends message to user

### calculate(expression)

Performs mathematical calculation

### retrieve\_policy(topic)

Returns relevant document chunks



## Tool Anatomy

```
def tool_name(parameter1: type, parameter2: type) -> return_type: """ Purpose: What this tool does  
When to use: Specific situations Input: Description  
of parameters Output: What it returns """ #  
Implementation return result
```

## ⌚ Key Principles

- Clear, descriptive names
- Well-defined input schemas
- Predictable output formats
- Good documentation



# Components: Memory



## Working Memory

Current task only

Steps taken so far, intermediate results, current context



## Session Memory

Current conversation

User preferences, context from earlier in chat, conversation flow



## Long-term Memory

Across sessions

Historical patterns, learned preferences, user profile



## Example Memory Structure

```
{ "user_id": "alice@company.com", "preferences": { "urgency_threshold": "high", "language": "en" }, "history": [ "requested PTO policy", "asked about benefits" ], "context": { "current_task": "schedule_meeting", "step": 2 } }
```



# System Instructions: The Agent's Constitution

"The core instructions that define behavior and boundaries"

## I Structure

### ROLE

Specific role and purpose

### GOAL

Primary objective

### CAPABILITIES

Available tools and their purposes

### CONSTRAINTS

Must do / Must not do

### PROCESS

Step-by-step workflow

### OUTPUT FORMAT

How to present results

## Example: Procurement Agent

### ROLE:

Procurement review assistant that validates vendor quotes

### GOAL:

Identify potential issues and flag items for human review

### CAPABILITIES:

- compare\_to\_history()
- check\_vendor\_status()
- calculate\_variance()

### CONSTRAINTS:

MUST flag quotes >\$10,000

MUST NOT auto-approve unapproved vendors

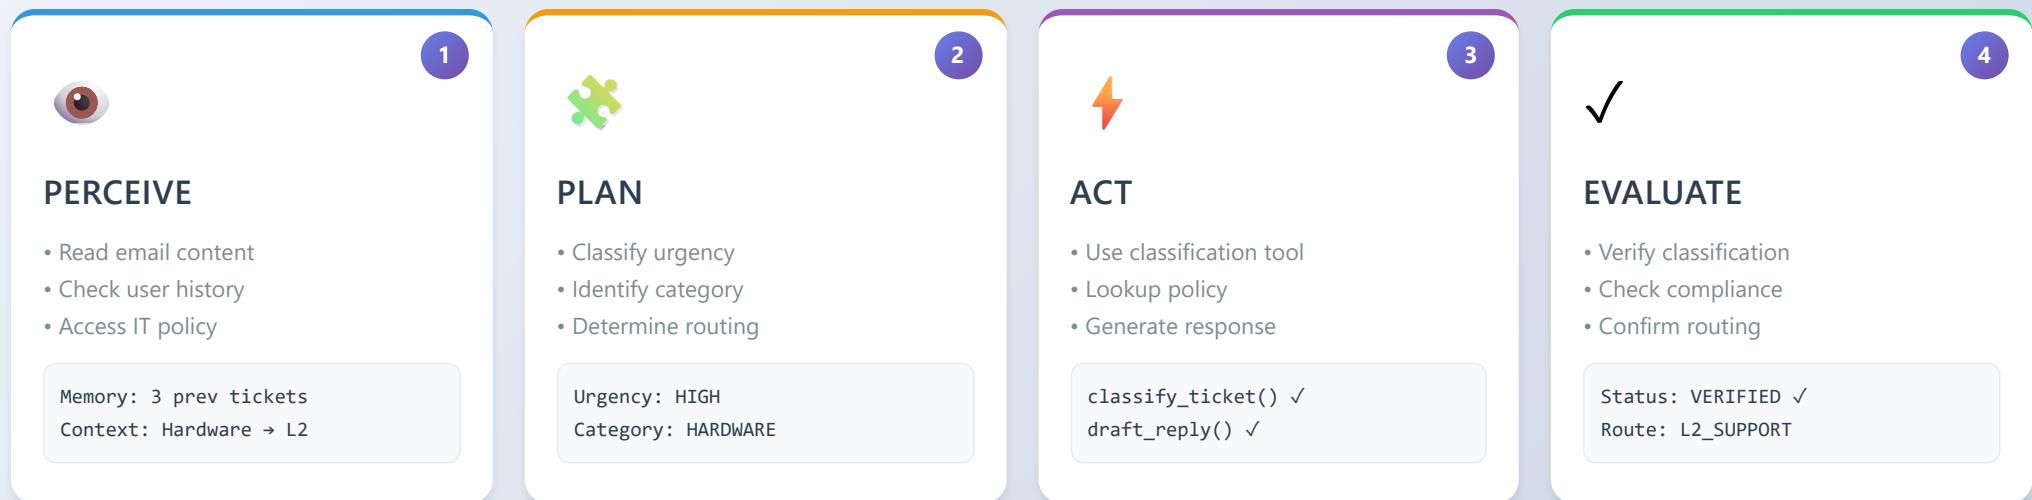
When data missing → "NEEDS REVIEW"



# Example: Email Triage Agent in Action

From: user@company.com

"My laptop won't turn on. URGENT!"



## Agent Output

TICKET STATUS

Escalated to L2

RESPONSE TIME

1.2 seconds

DRAFT REPLY

Auto-generated ✓

# Single vs Multi-Agent Systems



## Single Agent

*One agent, multiple tools*

### Agent



Tool 1



Tool 2



Tool 3

### Simple Architecture

Direct control flow, easier to debug

### Lower Latency

No coordination overhead

### Cost Effective

Single LLM call per decision



## Multi-Agent

*Specialized delegation system*

### Manager Agent



Research



Analysis



Writing

### Specialization

Each agent optimized for tasks

### Parallel Processing

Multiple tasks simultaneously

### Complex Workflows

Multi-step, branching processes



# Industry Use Cases



## HR Support Agent

Employee Services

Answers employee questions about policies, benefits, and procedures using company handbook.

- ✓ 24/7 availability
- ✓ Consistent answers
- ✓ Reduces HR workload by 60%



## Financial Analyst

Data Analysis

Analyzes financial data, generates reports, and flags anomalies for human review.

- ✓ Faster report generation
- ✓ Early anomaly detection
- ✓ Standardized formats



## Customer Service

Support Automation

Handles order inquiries, tracks shipments, and resolves common customer issues.

- ✓ Instant responses
- ✓ Multi-language support
- ✓ Escalates complex cases



## Legal Research

Document Analysis

Searches case law, summarizes precedents, and identifies relevant statutes.

- ✓ Hours saved per case
- ✓ Comprehensive coverage
- ✓ Citation accuracy



# Agent Canvas

"One-page business case for your agent"



## Problem & Value

- What problem?
- For whom?
- Why does it matter?



## Target User

- Who interacts?
- User context?
- Technical level?



## Success KPIs

- Measurable metrics
- How measured?
- Target values?



## Inputs & Outputs

- What comes in?
- What goes out?
- Format specs?



## Data Sources

- What data needed?
- Access method?
- Update frequency?



## Constraints

- Privacy requirements?
- Latency limits?
- Budget constraints?



# Agent Spec Sheet

"Technical blueprint for implementation"



## System Instruction

### ROLE

Clear definition of agent's purpose and identity

### GOAL

Primary objective the agent pursues

### CAPABILITIES

List of available tools and when to use them

### CONSTRAINTS

- MUST do / MUST NOT do
- Safety guardrails
- Fallback behaviors



## Implementation Details

### TOOLS

- Input/output schemas
- Example usage
- Error handling

### MEMORY

- What to store
- When to store
- Retention policy

### EVALUATION

- Key metrics
- Pass/fail thresholds
- Test cases (min 5)



# Testing Your Agent

## Happy Path

Standard scenarios where everything works as expected

"Employee asks about PTO policy → Agent returns correct policy section"

## Guardrails

Testing safety boundaries and constraints

"Request to access unauthorized data → Agent refuses appropriately"

## Edge Cases

Unusual inputs or boundary conditions

"Query with misspellings → Agent still understands intent"

## Error Handling

How agent behaves when things go wrong

"Database unavailable → Agent provides helpful error message"



## Best Practices

- ✓ Clear pass criteria for each test
- ✓ Document expected vs actual behavior
- ✓ Automate repetitive test cases

- ✓ Test with real user data
- ✓ Minimum 5 tests before deployment
- ✓ Include stakeholders in test review



# Risks & Guardrails



## Hallucination

Agent confidently states incorrect information

### Mitigation:

Require citations, verify against sources, confidence thresholds



## Privacy Leaks

Agent exposes sensitive or personal data

### Mitigation:

Access controls, data filtering, audit logs, PII detection



## Cost Overruns

Unexpected API calls or infinite loops

### Mitigation:

Rate limits, max iterations, budget alerts, caching



# Essential Guardrails



### Timeouts

Max execution time per task



### Iteration Limits

Max steps in agent loop



### Human-in-the-Loop

Required approval gates



### Audit Logging

Track all agent actions



### Scope Constraints

Clear boundaries of authority



### Output Validation

Verify format and content

# Workshop Activity

"Design Your Own Agent"

1

## Brainstorm

5 minutes

Identify a business problem in your organization that could benefit from an agent

2

## Design

10 minutes

Complete the Agent Canvas and draft your Agent Spec Sheet

3

## Test

5 minutes

Create 5 test cases covering happy path, edge cases, and guardrails



## Your Deliverables



### Agent Canvas

One-page business case



### Agent Spec

Technical blueprint



### 5 Tests

Success criteria



## Team Roles (Rotating Each Session)



### Product Owner

(PO)

Ensure technical choices align with business goals



### Agent Architect

(AA)

Lead architecture pattern selection and workflow design



### Toolsmith

(TS)

Implement tools and establish data connections



### Evaluator

(EV)

Build test workflow and track performance metrics

# ★ Key Takeaways ★

1



## Agents are goal-driven

They perceive, plan, and act autonomously—not just execute commands

2



## The Loop is Everything

Perceive → Plan → Act → Evaluate. Master this pattern

3



## Instructions = Constitution

System instructions encode business logic, safety rules, and boundaries

4



## Guardrails are Essential

Safety isn't optional. Build constraints and human gates from day one

5



## Test or Don't Trust

Define clear, measurable success criteria before deployment

6



## Start Small, Scale Smart

Better to do one thing really well than many things poorly



# Next Steps & Lecture 2 Preview

## Lecture 2: Patterns, Tools & Memory

*From Design → Implementation*



### Architecture Patterns

ReAct, Plan-and-Execute, Chain-of-Thought



### Implementing Tools

Build your first real agent tools



### Adding Memory

Context and state management



### Running Evaluations

Test and validate your agent

## Before Next Lecture

✓ Incorporate peer feedback into your Canvas and Spec

✓ Validate one key assumption (data availability, tool feasibility)

✓ Prepare questions about patterns and architecture