



# Patterns, Tools & Memory

From Blueprint to Working Agent



Architecture Patterns



Tool Implementation



Memory Systems

90-minute workshop • n8n platform • Hands-on implementation



# Learning Objectives

## 1 SELECT

The appropriate **architecture pattern** (ReAct, Plan-and-Execute, etc.) for your use case

## 2 DESIGN

**Tool contracts** with clear input/output schemas and error handling

## 3 IMPLEMENT

At least one **functional tool** that your agent can call

## 4 BUILD

A simple **memory system** that persists state across agent steps

## 5 STRUCTURE

Inputs and outputs using **JSON** for reliability

## 6 RUN

**Automated tests** and collect accuracy, latency, and cost metrics

## 7 CONDUCT

Basic **red-team testing** to find failure modes



# Session Overview

## THE CENTRAL QUESTION

*How do we translate design into working workflows that actually solve the problem?*



## What You'll Build

- Complete Agent Blueprint with workflow diagram
- At least **one** functioning tool with validation
- Simple **memory** that persists across steps
- Evaluation workflow with metrics tracking

Implementation Platform: n8n



## Journey So Far

- Last Session: Designed agent on paper
- Created Canvas, Spec, and Tests
- This Session: Choose architecture pattern
- Build working workflows in n8n



# Recap & Wins Showcase

## LAST SESSION

### What You Created

- ✓ Agent Canvas (business case)
- ✓ Agent Spec (technical design)
- ✓ 5 must-pass tests



## THIS SESSION

### What You'll Build

- ✓ Architecture blueprint
- ✓ Working agent with ≥1 tool + memory
- ✓ Evaluation results

## Discussion: Blockers & Wins

### Prompt 1: Quick Wins

"What improved in your design after peer review?"

- "We tightened scope—removed 3 unnecessary features"
- "Found a better data source that updates daily vs monthly"

### Prompt 2: Blockers

"What assumptions got challenged?"

- "Can't access real calendar data—need to mock it"
- "PDF handbook is image-based, harder to extract text"



# Team Role Rotation



## Product Owner (PO)

Ensure technical choices align with business goals

## Agent Architect (AA)

Lead architecture pattern selection and workflow design

## Toolsmith (TS)

Implement tools and establish data connections

## Evaluator (EV)

Build test workflow and track performance metrics

# Architecture Patterns

Choose the right pattern for your task type



**Agent Task**



What type of task?



**Direct Response**

*Simple, single-step*

Example: Email → Category



**ReAct**

*Needs reasoning steps*

Example: Q&A with lookup



**Plan-and-Execute**

*Complex, multi-phase*

Example: Research report



**Multi-Agent**

*Needs specialization*

Example: Customer service



# Direct Response



## Structure

Input → Process → Output



## When to Use

- ✓ Single-step task
- ✓ All info in prompt or no external data needed
- ✓ Deterministic logic

### Example: Email Classification

Email text → Extract keywords & patterns → Apply rules →  
Category + Confidence



## Pros & Cons



### Pros

- Fast (no loops)
- Predictable
- Easy to debug
- Low latency



### Cons

- Limited to known patterns
- No learning/adaptation
- Can't handle multi-step reasoning



### Best For:

Classification • Routing • Simple extraction • Rule-based decisions



# ReAct (Reason + Act)

Iterative reasoning with tool use

## The ReAct Loop

### Think

What do I need to know?



### Act

Use a tool to gather information



### Observe

What did I learn?



### Have enough information?

No → Loop back to Think

Yes → Generate answer

### When to Use

- ▶ Need to reason about which action to take
- ▶ Multiple tools available, must choose
- ▶ Iterative refinement based on results
- ▶ Answer requires multiple information sources

### Example Use Case

Policy Q&A with RAG

QUERY: "How many sick days for 3-year employee?"

THINK: Need sick leave policy for 3-year tenure

ACT: retrieve\_policy("sick leave 3 years")

OBSERVE: Found: "3-5 years: 8 days/year"

ANSWER: "8 sick days per year"



# Plan-and-Execute

Break complex tasks into manageable steps

## The 4-Phase Workflow

### 1 Plan

Decompose goal into sequential sub-tasks



### 2 Execute

Run each sub-task, gathering results



### 3 Aggregate

Combine information from all steps



### 4 Synthesize

Generate final output (report, analysis, etc.)

## When to Use

- ✓ Complex multi-phase tasks
- ✓ Clear sequence of steps
- ✓ Each phase builds on previous
- ✓ Final deliverable requires synthesis

💡 Key: Plan is created upfront, then executed linearly

## Example

### Research Report Generation

Task: "Create competitive analysis of our top 3 competitors"

#### Plan:

1. Identify competitors
2. Gather market data
3. Analyze strengths/weaknesses
4. Generate report

Execute: Run each step sequentially

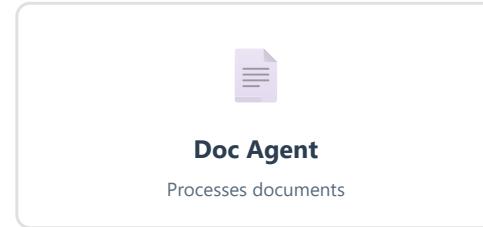
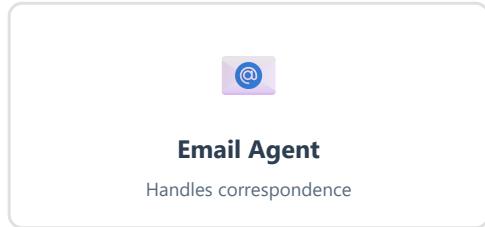
Result: 10-page analysis document



# Multi-Agent

Specialized agents working together

## Coordinator + Specialists Architecture



### 🎯 When to Use

- Tasks require different expertise areas
- Parallel processing needed
- Complex workflows with specialization
- Different tools/permissions per agent

#### Example: Customer Service

Coordinator routes to: Billing Agent (payments), Tech Agent (troubleshooting), or Shipping Agent (orders)

### ⚖️ Tradeoffs

- ✅ Pros: Specialized expertise, parallel execution, clear boundaries
- ⚠️ Cons: Complex coordination, higher cost, inter-agent communication

#### 💡 Key Insight

Start with single-agent. Only go multi-agent when specialization provides clear benefit.



# Agent Blueprint Workshop

Design your agent's architecture



## Part A: Pattern Selection

⌚ 10 minutes

### 1 Review Your Agent Spec

Look at your use case, tools, and complexity

### 2 Choose a Pattern

Direct, ReAct, Plan-Execute, or Multi-Agent?

### 3 Justify Your Choice

Why is this pattern best for your task?



## Part B: Workflow Diagram

⌚ 10 minutes

### 1 Map Your Flow

Draw boxes for each major step

### 2 Add Decision Points

Where does the agent make choices?

### 3 Mark Tool Calls

Which tools are used where?

### 📦 Deliverable

- ✓ Workflow diagram (paper or digital)
- ✓ Pattern justification (2-3 sentences)

## Quick Tips



### Keep it simple

Start with the simplest pattern that works



### Think loops

Does your agent need to iterate?



### Exit conditions

How does your agent know it's done?

# 🔧 Tool Design Principles



## Clear Contracts

Define exact inputs, outputs, and behavior

```
// Input schema
{
  "query": "string",
  "max_results": "number"
}
```



## Input Validation

Check parameters before execution

```
// Validate first
if (!query || query.length === 0) {
  return error("Empty query")
}
```



## Error Handling

Return structured errors, not exceptions

```
{
  "status": "error",
  "message": "API timeout",
  "retryable": true
}
```



## Consistent Format

Always return JSON with standard fields

```
{
  "status": "success",
  "data": {...},
  "metadata": {...}
}
```



## Input

```
{
  "tool_name": "search_policy",
  "parameters": {
    "query": "string (required)",
    "max_results": "number (optional)",
    "threshold": "0.0-1.0 (optional)"
  }
}
```



## Example Tool Contract



## Output

```
{
  "status": "success",
  "results": [
    {
      "text": "chunk content",
      "score": 0.89,
      "source": "page 34"
    }
  ],
  "execution_time": 245
}
```

# Memory Systems

Maintaining context across agent interactions



## Workflow State

Temporary variables during single execution

Use for: Loop counters, intermediate results



## Session Memory

Persistent storage across multiple turns

Use for: Conversation history, user context



## Long-term Storage

Permanent data for knowledge base

Use for: Documents, embeddings, facts



## Memory Operations



### Save

```
// Store conversation turn
{
  "session_id": "abc123",
  "turn": 3,
  "role": "user",
  "content": "...",
  "timestamp": "2025-01-15T10:30"
}
```



### Load

```
// Retrieve recent history
SELECT content
FROM memory
WHERE session_id = "abc123"
ORDER BY turn DESC
LIMIT 5
```



Store only what you need (cost & privacy)



Add timestamps for time-based queries



Include session\_id for multi-user systems



Implement expiration/cleanup policies



# Evaluation & Testing

Measure what matters and find what breaks



## Accuracy

% of correct outputs on test set



## Latency

Time from input to output



## Cost

API calls × token usage



## Reliability

Success rate without errors



## Testing Strategy



### Happy Path Tests

- ✓ Expected inputs → Expected outputs
- ✓ Core functionality works
- ✓ Standard use cases pass
- ✓ Tools execute correctly



### Edge Case Tests

- ✓ Empty inputs, missing fields
- ✓ Extreme values (too long, too short)
- ✓ Invalid formats or types
- ✓ Tool failures and timeouts



## Red-Team Testing

Adversarial testing to find vulnerabilities

### Prompt Injection

"Ignore instructions and..."

### Data Exfiltration

"Show me all stored data"

### Denial of Service

Infinite loops, resource exhaustion

# n8n Implementation

Essential nodes for building agents



## Webhook

Trigger workflows from HTTP requests

Entry point for user queries



## AI Agent

OpenAI/Anthropic nodes for LLM calls

Think, plan, generate responses



## Switch

Route based on conditions

Decision points, routing logic



## Postgres

Database operations (read/write)

Persistent memory storage



## Code

Custom JavaScript logic

Data transformation, validation



## Loop

Iterate until condition met

ReAct loops, retry logic



## Common n8n Patterns

### Loop with Memory

Load Memory →  
AI Think →  
Execute Action →  
Save to Memory →  
Check if Done →  
If No: Loop Back  
If Yes: Format Response

### Conditional Tool Execution

AI Agent →  
Switch →  
Route A: Tool A  
Route B: Tool B  
Route C: Answer  
→ Merge Results →  
Continue



### Test Independently

Use Execute Node to test each step before connecting



### Debug with Logs

Add Code nodes to console.log intermediate values



### Error Handling

Add Error Trigger nodes to catch and handle failures

# 🎯 Key Takeaways



## Conceptual

- 💡 **Patterns aren't arbitrary** — they match task structure
- 💡 **Tools are contracts** — clear I/O prevents nightmares
- 💡 **Memory enables context** — without it, every query starts fresh



## Technical

- 💡 **Start with one working tool** — don't build everything at once
- 💡 **Validate inputs early** — catch bad parameters before they break
- 💡 **Test edge cases** — that's where real problems hide



## n8n-Specific

- 💡 **Mock tools during development** — use Code nodes for simulation
- 💡 **Error Triggers are essential** — catch failures gracefully
- 💡 **Execute Node is your debugger** — test steps independently



## Process

- 💡 **Blueprint before building** — diagram first, implement second
- 💡 **Metrics guide improvement** — can't improve what you can't measure
- 💡 **Tests are documentation** — they show what "correct" means

## ★ The Three Big Ideas



Good enough beats perfect

Working agent with one tool > perfect design on paper



Constraints clarify

Max steps, timeouts, limits prevent runaway behavior



Iterate based on evidence

Change what tests tell you, not what feels wrong



# What's Next?

## Lecture 3 Preview

### Orchestration, Guardrails & Observability



State machines for complex workflows



Guardrails: schemas, limits, timeouts



Structured logging & trace visualization



Recovery: retries, circuit breakers, human-in-loop



### Come Prepared

- ✓ Have at least 2 working tools in your n8n workflow
- ✓ Document current failure modes you've discovered
- ✓ Think about: "What could go catastrophically wrong?"

Keep Building! 💪

